

CMPE 362

Digital Image Processing

Spatial Domain Operations II: Spatial Filtering

Asst. Prof. Dr. Aslı Gençtav
Department of Computer Engineering
TED University

Spatial domain operations

$$g(x, y) = T[f(x, y)]$$

$f(x, y)$: input image

$g(x, y)$: output image

T : operator defined over a neighborhood of (x, y)

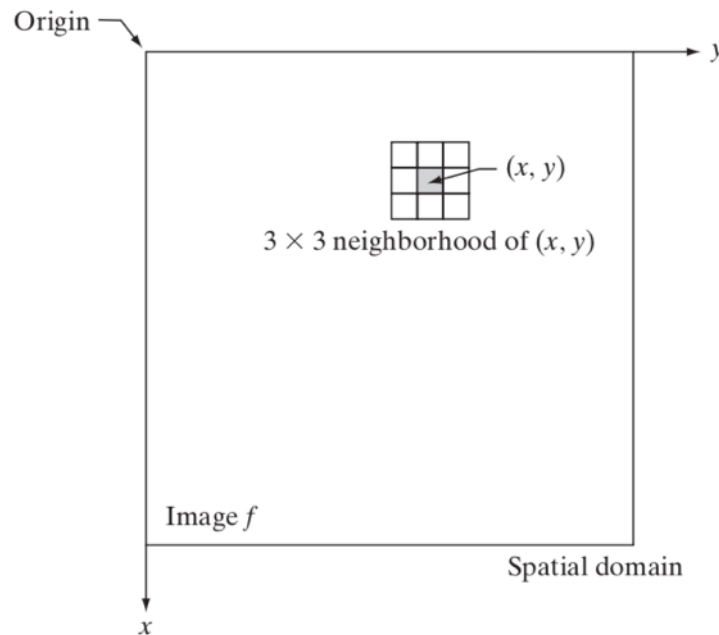


FIGURE 3.1

A 3×3 neighborhood about a point (x, y) in an image in the spatial domain. The neighborhood is moved from pixel to pixel in the image to generate an output image.

Spatial filtering


- Spatial filtering consists of
 1. a neighborhood (typically a small rectangle)
 2. a predefined operation
 - If the operation is linear, the filter is linear.
 - Otherwise, the filter is nonlinear.
- Uses
 - Enhance images
 - Noise reduction, resize, increase contrast, artistic effect, etc.
 - Extract information from images
 - Texture, edges, distinctive points, etc.
 - Detect patterns
 - Template matching

Linear spatial filtering

- Linear spatial filtering of an $M \times N$ image with a $m \times n$ filter is given by

$$g(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x + s, y + t)$$

where $m = 2a + 1$ and $n = 2b + 1$.



$$g(x, y) = w(-1, -1)f(x - 1, y - 1) + w(-1, 0)f(x - 1, y) + \dots \\ + w(0, 0)f(x, y) + \dots + w(1, 1)f(x + 1, y + 1)$$

input image f

| | 0 | 1 | 2 | 3 | 4 |
|---|----|----|----|----|----|
| 0 | 10 | 80 | 60 | 50 | 20 |
| 1 | 20 | 40 | 40 | 30 | 20 |
| 2 | 30 | 10 | 10 | 20 | 60 |
| 3 | 40 | 10 | 20 | 10 | 40 |
| 4 | 60 | 10 | 30 | 10 | 30 |

filter w

| | | |
|---|---|---|
| 1 | 2 | 1 |
| 2 | 4 | 2 |
| 1 | 2 | 1 |

output image g

| | 0 | 1 | 2 | 3 | 4 |
|---|---|---|---|---|---|
| 0 | | | | | |
| 1 | | ? | | | |
| 2 | | | | | |
| 3 | | | | | |
| 4 | | | | | |

$g(1, 1) = ?$

Poll 1

input image f

| | 0 | 1 | 2 | 3 | 4 |
|---|----|----|----|----|----|
| 0 | 10 | 80 | 60 | 50 | 20 |
| 1 | 20 | 40 | 40 | 30 | 20 |
| 2 | 30 | 10 | 10 | 20 | 60 |
| 3 | 40 | 10 | 20 | 10 | 40 |
| 4 | 60 | 10 | 30 | 10 | 30 |

filter w

| | | |
|----|----|----|
| 1 | 2 | 1 |
| 0 | 0 | 0 |
| -1 | -2 | -1 |

output image g

| | 0 | 1 | 2 | 3 | 4 |
|---|---|---|---|---|---|
| 0 | | | | | |
| 1 | | | ? | | |
| 2 | | | | | |
| 3 | | | | | |
| 4 | | | | | |

$$g(1, 2) = ?$$

Poll 2

Linear spatial filtering

Correlation vs. Convolution

- **Correlation:**

1. Move the filter mask to a location
2. Compute the sum of products
3. Go to 1

- **Convolution:**

1. Rotate the filter by 180 degrees
(flip the filter in both dimensions: bottom to top, right to left)
2. Correlation

Linear spatial filtering


Correlation vs. Convolution

- **Correlation:**

1. Move the filter mask to a location
2. Compute the sum of products
3. Go to 1

- **Convolution:**

1. Rotate the filter by 180 degrees
(flip the filter in both dimensions: bottom to top, right to left)
2. Correlation



Linear spatial filtering

Correlation vs. Convolution

- Correlation:

$$(w \star f)(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x + s, y + t) \quad (3-34)$$

- Convolution:

$$(w \star f)(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x - s, y - t) \quad (3-35)$$

Correlation

(a) ↖ Origin f w
 0 0 0 1 0 0 0 0 1 2 4 2 8

(b) ↓
 0 0 0 1 0 0 0 0
 1 2 4 2 8
 ↑ Starting position alignment

(c) — Zero padding —
 ↙ ↘
 0 0 0 0 0 1 0 0 0 0 0 0
 1 2 4 2 8
 ↑ Starting position

(d) 0 0 0 0 0 1 0 0 0 0 0 0
 1 2 4 2 8
 ↑ Position after 1 shift

(e) 0 0 0 0 0 1 0 0 0 0 0 0
 1 2 4 2 8
 ↑ Position after 3 shifts

(f) 0 0 0 0 0 1 0 0 0 0 0 0
 1 2 4 2 8
 Final position —↑


Correlation result

(g) 0 8 2 4 2 1 0 0

Extended (full) correlation result

(h) 0 0 0 8 2 4 2 1 0 0 0 0

Correlation




Correlation result

(g) 0 8 2 4 2 1 0 0

Extended (full) correlation result

(h) 0 0 0 8 2 4 2 1 0 0 0 0

Convolution



Convolution result

(o) 0 1 2 4 2 8 0 0

Extended (full) convolution result

(p) 0 0 0 1 2 4 2 8 0 0 0 0

Diagram illustrating the padding process for a 1D convolution:

(a) Original input f (size 5) and kernel w (size 3). The input f is shown as a 5x5 grid of zeros, with the third row containing the values 1, 2, 3. The kernel w is shown as a 3x3 grid of zeros, with the third row containing the values 4, 5, 6.

(b) Padded input f (size 7). The input f is shown as a 7x7 grid of zeros, with the third row containing the values 1, 2, 3. The kernel w is shown as a 3x3 grid of zeros, with the third row containing the values 4, 5, 6. The padding is 1 on each side, making the total size 7.

Correlation

Convolution

TABLE 3.5
 Some fundamental properties of convolution and correlation. A dash means that the property does not hold.

| Property | Convolution | Correlation |
|--------------|---|---|
| Commutative | $f \star g = g \star f$ | — |
| Associative | $f \star (g \star h) = (f \star g) \star h$ | — |
| Distributive | $f \star (g + h) = (f \star g) + (f \star h)$ | $f \star (g + h) = (f \star g) + (f \star h)$ |

Linear spatial filtering

Correlation vs. Convolution

- What if the values of the filter are symmetric about its center?

Practical matters

What about near the edge?

- the filter window falls off the edge of the image
- need to extrapolate
- methods:
 - clip filter (black) 000000|abcdefgh|000000
`g = cv2.filter2D(f, -1, w, borderType=cv2.BORDER_CONSTANT)`
 - copy edge aaaaaa|abcdefgh|hhhhhhh
`g = cv2.filter2D(f, -1, w, borderType=cv2.BORDER_REPLICATE)`
 - reflect across edge fedcba|abcdefgh|hgfedcb
`g = cv2.filter2D(f, -1, w, borderType=cv2.BORDER_REFLECT)`
 - reflect across edge gfedcb|abcdefgh|gfedcba
`g = cv2.filter2D(f, -1, w, borderType=cv2.BORDER_REFLECT_101)`

Poll 4

| | 0 | 1 | 2 | 3 |
|---|----|----|----|----|
| 0 | 10 | 80 | 60 | 50 |
| 1 | 20 | 40 | 40 | 30 |
| 2 | 30 | 10 | 10 | 20 |
| 3 | 40 | 10 | 20 | 10 |

input image f

| | 0 | 1 | 2 |
|---|---|---|---|
| 0 | 1 | 1 | 1 |
| 1 | 1 | 2 | 1 |
| 2 | 2 | 4 | 2 |
| 3 | 1 | 2 | 1 |
| 4 | 1 | 1 | 1 |

filter w

BORDER_CONSTANT
BORDER_REPLICATE
BORDER_REFLECT
BORDER_REFLECT_101

A

| | 0 | 1 | 2 | 3 |
|---|-----|-----|-----|-----|
| 0 | 700 | 840 | 880 | 800 |
| 1 | | | | |
| 2 | | | | |
| 3 | | | | |

B

| | 0 | 1 | 2 | 3 |
|---|-----|------|------|-----|
| 0 | 600 | 1030 | 1130 | 970 |
| 1 | | | | |
| 2 | | | | |
| 3 | | | | |

C

| | 0 | 1 | 2 | 3 |
|---|-----|-----|------|-----|
| 0 | 580 | 980 | 1050 | 910 |
| 1 | | | | |
| 2 | | | | |
| 3 | | | | |

D

| | 0 | 1 | 2 | 3 |
|---|-----|-----|-----|-----|
| 0 | 320 | 650 | 690 | 450 |
| 1 | | | | |
| 2 | | | | |
| 3 | | | | |

Smoothing spatial filters

- Linear filters
 - Averaging filter (Box filter)
 - Gaussian filter
- Nonlinear filters
 - Median filter

Common types of noise

Poll 5

- Salt and pepper noise:



Original



Salt and pepper noise

- Impulse noise:



Impulse noise

- Gaussian noise:



Gaussian noise

Common types of noise

- **Salt and pepper noise:** random occurrences of black and white pixels
- **Impulse noise:** random occurrences of white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian normal distribution



Original



Salt and pepper noise






Impulse noise



Gaussian noise

Gaussian noise



$$f(x, y) = \overbrace{\hat{f}(x, y)}^{\text{Ideal Image}} + \overbrace{\eta(x, y)}^{\text{Noise process}}$$

Gaussian i.i.d. ("white") noise:
 $\eta(x, y) \sim \mathcal{N}(\mu, \sigma)$

Averaging filter

$F[x, y]$

| | | | | | | | | | |
|---|---|----|----|----|----|----|----|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

$G[x, y]$

| | | | | | | | | | |
|--|---|----|--|--|--|--|--|--|--|
| | | | | | | | | | |
| | 0 | 10 | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |

Averaging filter

$F[x, y]$

| | | | | | | | | | |
|---|---|----|----|----|----|----|----|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

$G[x, y]$

| | | | | | | | | | |
|--|---|----|----|--|--|--|--|--|--|
| | | | | | | | | | |
| | 0 | 10 | 20 | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |

Averaging filter

$$F[x, y]$$

| | | | | | | | | | |
|---|---|----|----|----|----|----|----|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

$$G[x, y]$$

| | | | | | | | | | |
|--|---|----|----|----|--|--|--|--|--|
| | | | | | | | | | |
| | 0 | 10 | 20 | 30 | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |

Averaging filter

$$F[x, y]$$

| | | | | | | | | | |
|---|---|----|----|----|----|----|----|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

$$G[x, y]$$

| | | | | | | | | | |
|--|---|----|----|----|----|--|--|--|--|
| | | | | | | | | | |
| | 0 | 10 | 20 | 30 | 30 | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |

Averaging filter

$$F[x, y]$$


| | | | | | | | | | |
|---|---|----|----|----|----|----|----|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

$$G[x, y]$$

| | | | | | | | | | |
|--|----|----|----|----|----|----|----|----|--|
| | | | | | | | | | |
| | 0 | 10 | 20 | 30 | 30 | 30 | 20 | 10 | |
| | 0 | 20 | 40 | 60 | 60 | 60 | 40 | 20 | |
| | 0 | 30 | 60 | 90 | 90 | 90 | 60 | 30 | |
| | 0 | 30 | 50 | 80 | 80 | 90 | 60 | 30 | |
| | 0 | 30 | 50 | 80 | 80 | 90 | 60 | 30 | |
| | 0 | 20 | 30 | 50 | 50 | 60 | 40 | 20 | |
| | 10 | 20 | 30 | 30 | 30 | 30 | 20 | 10 | |
| | 10 | 10 | 10 | 0 | 0 | 0 | 0 | 0 | |
| | | | | | | | | | |


Averaging filter

- What values belong in the kernel H for the moving average example?



$$G = H \otimes F$$

Smoothing by averaging filter




box filter

white = high value, black = low value



original



filtered

What if the filter size increases?

Gaussian filter

A Gaussian kernel gives less weight to pixels further from the center of the filter.

| | | | | | | | | | |
|---|---|----|----|----|----|----|----|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

$F[x, y]$


| | | |
|---|---|---|
| 1 | 2 | 1 |
| 2 | 4 | 2 |
| 1 | 2 | 1 |

$\frac{1}{16}$

$H[u, v]$

This kernel is an approximation of a Gaussian function:

$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$

$$\sigma = 1$$

$$h(u, v) = \frac{1}{2\pi} e^{-\frac{u^2+v^2}{2}}$$

$$h(0,0) = \frac{1}{2\pi} e^{-\frac{0^2+0^2}{2}} = 0.1592$$

$$h(-1,0) = \frac{1}{2\pi} e^{-\frac{(-1)^2+0^2}{2}} = 0.0965$$

$$h(0,-1) = h(0,1) = h(1,0) = h(-1,0)$$

$$h(-1,-1) = \frac{1}{2\pi} e^{-\frac{(-1)^2+(-1)^2}{2}} = 0.0585$$



$$h(-1,1) = h(1,-1) = h(1,1) = h(-1,-1)$$

| | -1 | 0 | 1 |
|----|--------|--------|--------|
| -1 | 0.0585 | 0.0965 | 0.0585 |
| 0 | 0.0965 | 0.1592 | 0.0965 |
| 1 | 0.0585 | 0.0965 | 0.0585 |

Normalized to have sum 1


| | -1 | 0 | 1 |
|----|--------|--------|--------|
| -1 | 0.0751 | 0.1238 | 0.0751 |
| 0 | 0.1238 | 0.2043 | 0.1238 |
| 1 | 0.0751 | 0.1238 | 0.0751 |

Averaging filter vs. Gaussian filter




Gaussian filters

- What parameters matter here?
- **Size** of kernel or mask
 - Note, Gaussian function has infinite support, but discrete filters use finite kernels




$\sigma = 5$
with 10 x 10 kernel




$\sigma = 5$
with 30 x 30 kernel

Gaussian filters

- What parameters matter here?
- **Variance** of Gaussian: determines extent of smoothing



$\sigma = 2$
with 30 x 30 kernel




$\sigma = 5$
with 30 x 30 kernel

Gaussian filters



Choosing kernel width

- Rule of thumb: set filter half-width to about 3σ




Smoothing with a Gaussian


Parameter σ is the scale/width/spread of the Gaussian kernel, and controls the amount of smoothing.




...





Gaussian filters




$$\sigma = 1$$



$$\sigma = 5$$



$$\sigma = 10$$



$$\sigma = 30$$

Gaussian Kernel

- Convolution with self is another Gaussian
 - So can smooth with small- σ kernel, repeat, and get same result as larger- σ kernel would have
 - Convoluting two times with Gaussian kernel with std. dev. σ is same as convoluting once with kernel with std. dev. $\sigma\sqrt{2}$

Gaussian Kernel

- Convolution with self is another Gaussian
 - So can smooth with small- σ kernel, repeat, and get same result as larger- σ kernel would have
 - Convoluting two times with Gaussian kernel with std. dev. σ is same as convoluting once with kernel with std. dev. $\sigma\sqrt{2}$
- *Separable* kernel
 - Factors into product of two 1D Gaussians
 - Discrete example:
$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

Separability of Gaussian filter

$$\begin{aligned} G_{\sigma}(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y

In this case, the two functions are the (identical) 1D Gaussian

Separability

- In some cases, filter is separable, and we can factor into two steps:
 - Convolve all rows
 - Convolve all columns

Separability example

2D convolution
(center location only)

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array} = \begin{array}{l} = 2 + 6 + 3 = 11 \\ = 6 + 20 + 10 = 36 \\ = 4 + 8 + 6 = 18 \\ \hline 65 \end{array}$$

The filter factors into
a product of 1D
filters:

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array}$$

Perform convolution
along rows:

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array} \begin{array}{|c|c|c|} \hline 2 & 3 & 3 \\ \hline 3 & 5 & 5 \\ \hline 4 & 4 & 6 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline & 11 & \\ \hline & 18 & \\ \hline & 18 & \\ \hline \end{array}$$

Followed by
convolution along the
remaining column:

$$\begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} \begin{array}{|c|c|c|} \hline & 11 & \\ \hline & 18 & \\ \hline & 18 & \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline & & \\ \hline & 65 & \\ \hline & & \\ \hline \end{array}$$


Why is separability useful?


- What is the complexity of filtering an $N \times N$ image with an $M \times M$ kernel?
- What if the kernel is separable?


Why is separability useful?

- What is the complexity of filtering an $N \times N$ image with an $M \times M$ kernel?
 - $O(N^2 M^2)$
- What if the kernel is separable?
 - $O(N^2 M)$

Predict outputs using correlation filtering


$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} = ?$$


$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix} = ?$$


$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} = ?$$



Input image

A

| | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 1 | 0 |
| 0 | 0 | 0 |

B

| | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 0 | 1 |
| 0 | 0 | 0 |


C

| | | | |
|---------------|---|---|---|
| $\frac{1}{9}$ | 1 | 1 | 1 |
| | 1 | 1 | 1 |
| | 1 | 1 | 1 |

D


| | | | | | | |
|---|---|---|----------------|---|---|---|
| 0 | 0 | 0 | $-\frac{1}{9}$ | 1 | 1 | 1 |
| 0 | 2 | 0 | | 1 | 1 | 1 |
| 0 | 0 | 0 | | 1 | 1 | 1 |

1



Shifted *left* by 1 pixel

2



Smoothed

3



Sharpened

4



No change

Practice with linear filters



Original

| | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 1 | 0 |
| 0 | 0 | 0 |

?

Practice with linear filters



Original

| | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 1 | 0 |
| 0 | 0 | 0 |



Filtered
(no change)

Practice with linear filters



Original

| | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 0 | 1 |
| 0 | 0 | 0 |

?

Practice with linear filters



Original

| | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 0 | 1 |
| 0 | 0 | 0 |



Shifted *left*
by 1 pixel

Practice with linear filters



Original

 $\frac{1}{9}$

| | | |
|---|---|---|
| 1 | 1 | 1 |
| 1 | 1 | 1 |
| 1 | 1 | 1 |

?


Practice with linear filters



Original

$$\frac{1}{9}$$

| | | |
|---|---|---|
| 1 | 1 | 1 |
| 1 | 1 | 1 |
| 1 | 1 | 1 |



Blur (with a
box filter)

Practice with linear filters



Original

| | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 2 | 0 |
| 0 | 0 | 0 |

—

$\frac{1}{9}$

| | | |
|---|---|---|
| 1 | 1 | 1 |
| 1 | 1 | 1 |
| 1 | 1 | 1 |

?

(Note that filter sums to 1)

Practice with linear filters



Original

| | | |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 2 | 0 |
| 0 | 0 | 0 |

—

$\frac{1}{9}$


| | | |
|---|---|---|
| 1 | 1 | 1 |
| 1 | 1 | 1 |
| 1 | 1 | 1 |




Sharpening filter:

Accentuates differences with local average

Filtering examples: sharpening



before



after

Sharpening

What does blurring take away?



-




=



Let's add it back:



+



=




Unsharp Masking

- Blur the original image
- Subtract the blurred one from the original one
- Add the mask to the original

$$g(x, y) = f(x, y) + k(f(x, y) - \overline{f(x, y)})$$

$k = 1$ unsharp masking

$k > 1$ highboost filtering



Unsharp Masking




original



sharpened

Unsharp Masking





input image

| | | |
|----|----|----|
| 1 | 2 | 1 |
| 0 | 0 | 0 |
| -1 | -2 | -1 |


Sobel

A


| | | |
|---|---|----|
| 1 | 0 | -1 |
| 2 | 0 | -2 |
| 1 | 0 | -1 |

Sobel

B




**1 vertical edge
(absolute value)**



**2 horizontal edge
(absolute value)**


Other filters



| | | |
|---|---|----|
| 1 | 0 | -1 |
| 2 | 0 | -2 |
| 1 | 0 | -1 |

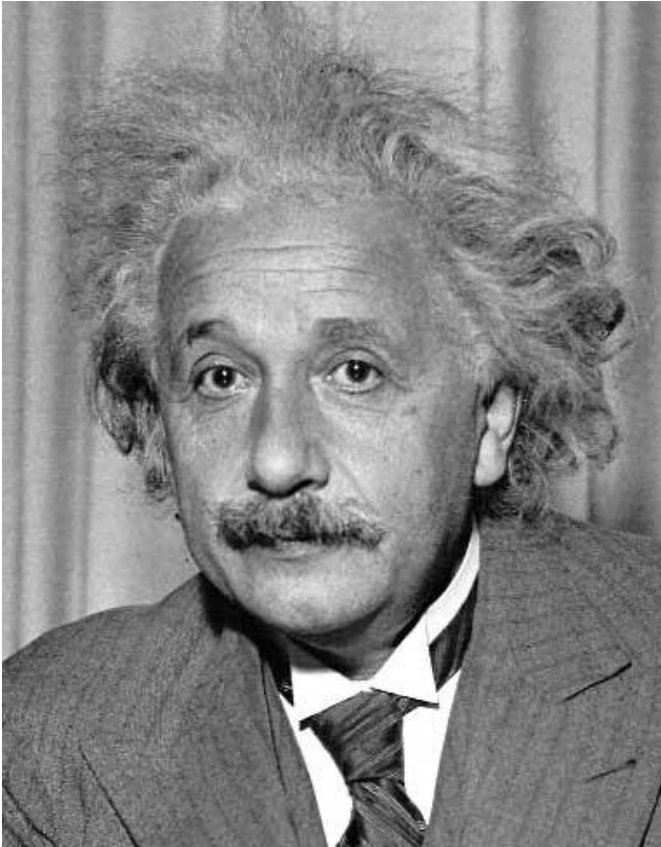
Sobel

B



vertical edge
(absolute value)


Other filters



| | | |
|----|----|----|
| 1 | 2 | 1 |
| 0 | 0 | 0 |
| -1 | -2 | -1 |

Sobel

A




horizontal edge
(absolute value)

Nonlinear Filters

- Median filter

Median filter

- A **median filter** operates over a window by selecting the median intensity in the window



- Is median filtering linear?

| | 0 | 1 | 2 | 3 |
|---|----|----|----|----|
| 0 | 10 | 80 | 60 | 50 |
| 1 | 20 | 40 | 40 | 30 |
| 2 | 30 | 10 | 10 | 20 |
| 3 | 40 | 10 | 20 | 10 |

input image f

3x3
median filter


| | 0 | 1 | 2 | 3 |
|---|---|---|---|---|
| 0 | | | | |
| 1 | | | | |
| 2 | | ? | | |
| 3 | | | | |

output image g



$$g(2, 1) = ?$$

Median filter

Salt-and-pepper noise



Median filtered




Plots of a row of the image


Median filter

- What advantage does median filtering have over Gaussian filtering?
 - Robustness to outliers
 - Median filter is edge preserving

filters have width 5 :



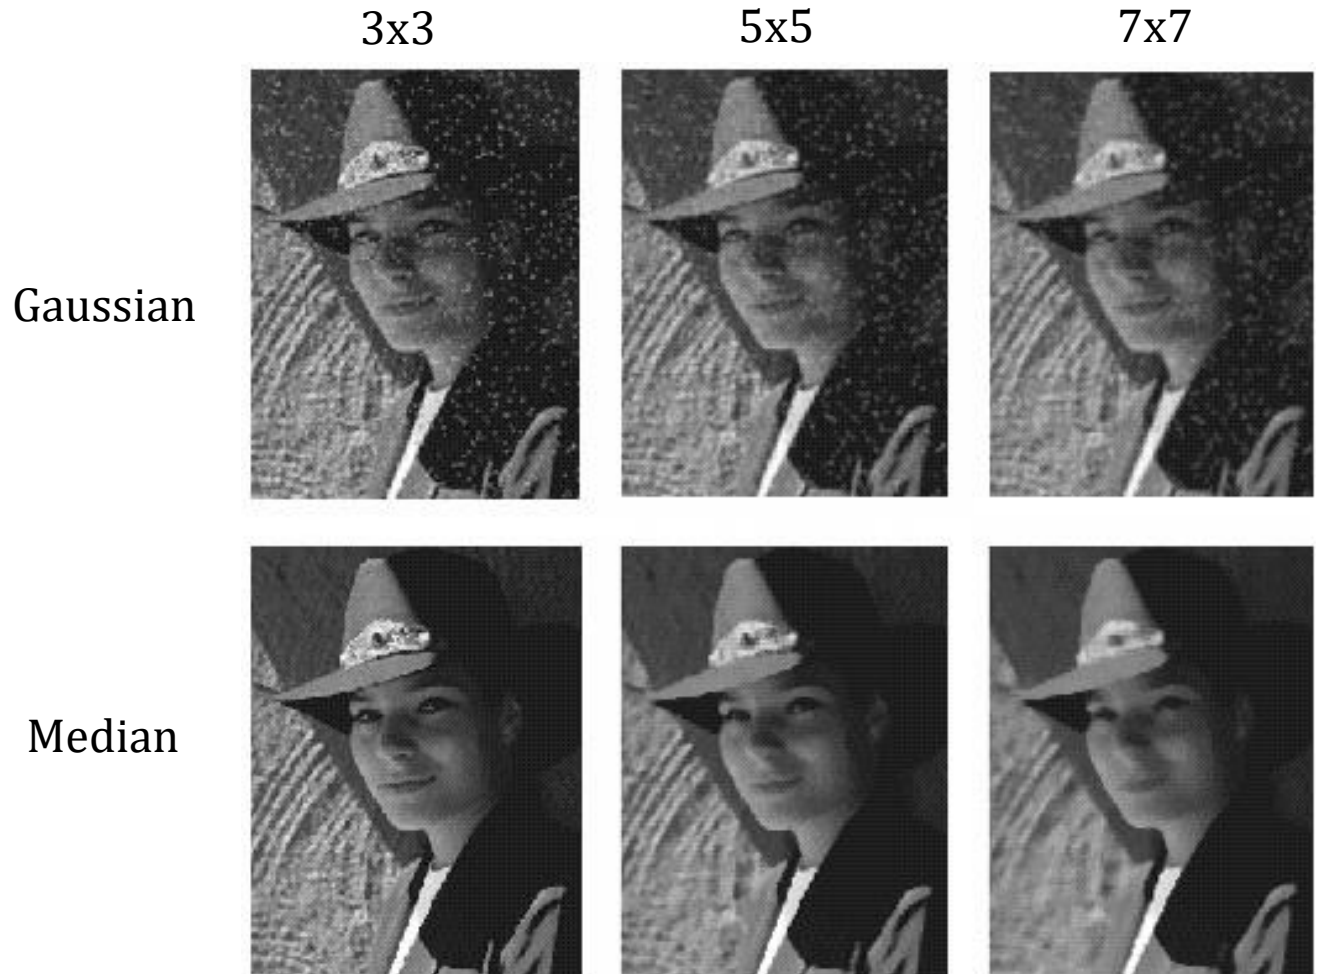
Averaging filter vs. Median filter



a b c

FIGURE 3.35 (a) X-ray image of circuit board corrupted by salt-and-pepper noise. (b) Noise reduction with a 3×3 averaging mask. (c) Noise reduction with a 3×3 median filter. (Original image courtesy of Mr. Joseph E. Pascente, Lixi, Inc.)

Gaussian filter vs. Median filter



Isolated clusters of pixels that are light or dark with respect to their neighbors, and whose area is less than $m^2/2$ (one-half the filter area), are forced by an $m \times m$ median filter to have the value of the median intensity of the pixels in the neighborhood.

Week 04 – Hands on activity

- Prepare and submit a Jupyter Notebook file containing the code and the results for the following Task

Task

- Read a colored image of your choice.
- Convert it into a gray-scale image and display the result.
- Filter the gray-scale image using each of the following filters.

$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

- Display the result using a colormap other than gray.
- Display the absolute value of the result using a colormap that you choose.