



CMPE 362

Digital Image Processing

Fitting: Hough Transform

Fitting

- Associate a model with observed features (such as edges)



[Fig from Marszalek & Schmid, 2007]


- For example, the model could be a line, a circle, or an arbitrary shape.

Fitting: Main idea

- Choose a parametric model to represent a set of features such as edges
- Membership criterion is not local
 - We can not tell whether a point belongs to a given model just by looking at that point
- Three main questions:
 - What model represents this set of features best?
 - Which of several model instances gets which feature?
 - How many model instances are there?
- Computational complexity is important
 - It is infeasible to examine every possible set of parameters and every possible combination of features



Case study: Line fitting

- Why fit lines?
Many objects could be characterized by presence of straight lines.



- Wait, why aren't we done just by running edge detection?

Difficulties of line fitting



- **Extra** edge points (clutter), multiple models:
 - which points go with which line, if any?
- Only some parts of each line are detected, and some parts are **missing**:
 - how to find a line that bridges missing evidence?
- **Noise** in measured edge points, orientations:
 - how to detect true underlying parameters?

Voting



- It's not feasible to check all combinations of features by fitting a model to each possible subset.
- Voting is a general technique where we let each feature vote for all models that are compatible with it.
 - Cycle through features, cast votes for model parameters.
 - Look for model parameters that receive a lot of votes.
- Noise & clutter features will cast votes too, but typically their votes should be inconsistent with the majority of “good” features.

Fitting lines: Hough transform


- Given points that belong to a line, what is the line?
- How many lines are there?
- Which points belong to which lines?
- Hough Transform is a voting technique that can be used to answer all of these questions.

Main idea:

1. Record vote for each possible line on which each edge point lies.
2. Look for lines that get many votes.




Finding lines in an image: Hough space



Connection between image (x, y) and Hough (m, b) spaces

- A line in the image corresponds to a point in Hough space


Finding lines in an image: Hough space



Connection between image (x, y) and Hough (m, b) spaces

- A line in the image corresponds to a point in Hough space
- To go from image space to Hough space:
 - given a set of points (x, y) , find all (m, b) such that $y = mx + b$
- **What does a point (x_0, y_0) in the image space map to?**
 - Answer: the solutions of $b = -x_0m + y_0$
 - this is a line in Hough space

Finding lines in an image: Hough space



What are the line parameters for the line that contains both (x_0, y_0) and (x_1, y_1) ?

- It is the intersection of the lines $b = -x_0m + y_0$ and $b = -x_1m + y_1$

Finding lines in an image: Hough space





image space




Hough (parameter) space

How can we use this to find the most likely parameters (m, b) for the most prominent line in the image space?

- Let each edge point in image space *vote* for a set of possible parameters in Hough space
- Accumulate votes in discrete set of bins; parameters with the most votes indicate the most prominent lines in image space.

Polar representation for lines

Issues with usual (m, b) parameter space: can take on infinite values, undefined for vertical lines.



d : perpendicular distance to line from origin

θ : angle the perpendicular makes with the x-axis

$$x \cos \theta - y \sin \theta = d$$

Point in image space \rightarrow sinusoid segment in Hough space

Hough transform algorithm


Using the polar parameterization:

$$x \cos \theta - y \sin \theta = d$$

Basic Hough transform algorithm


1. Initialize $H[d, \theta] = 0$
2. for each edge point $I[x, y]$ in the image
 for $\theta = [\theta_{\min} \text{ to } \theta_{\max}]$ // some quantization
 $d = x \cos \theta - y \sin \theta$
 $H[d, \theta] += 1$
3. Find the value or values of (d, θ) where $H[d, \theta]$ is maximum
4. The detected line in the image is given by
 $d = x \cos \theta - y \sin \theta$

H: accumulator array (votes)





Example: What was the shape?

Square



Example: Hough transform for straight lines




Which line generated this peak?

Original image




Canny edges



Decode the
vote space.

d



Impact of noise on Hough





Image space
edge coordinates



Votes

What difficulty does this present for an implementation?

Impact of noise on Hough





Image space
edge coordinates



θ
Votes

Here, everything appears to be “noise”, or random edge points, but we still see peaks in the vote space.

Voting: practical tips

- Minimize irrelevant tokens first
- Choose a good grid / discretization

← Too fine ? Too coarse →


- Vote for neighbors, also (smoothing in accumulator array)
- Use direction of edge to reduce parameters by 1
- To read back which points voted for “winning” peaks, keep tags on the votes.

Hough transform for circles

- Circle: center (a, b) and radius r

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For a fixed radius r , unknown gradient direction




Hough transform for circles

- Circle: center (a, b) and radius r

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For an unknown radius r , unknown gradient direction



Hough transform for circles

- Circle: center (a, b) and radius r

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For an unknown radius r , unknown gradient direction






Image space




Hough space

Hough transform for circles

- Circle: center (a, b) and radius r

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For an unknown radius r , **known** gradient direction



Hough transform for circles

For every edge pixel (x,y) :

For each possible radius value r :

For each possible gradient direction θ :

// or use estimated gradient at (x,y)

$a = x + r \cos(\theta)$ *// column*

$b = y - r \sin(\theta)$ *// row*

$H[a, b, r] += 1$


end

end


end

Example: detecting circles with Hough


Original



Edges




Votes: Penny




Example: detecting circles with Hough


Original



Edges



Votes: Quarter



Note: a different Hough transform (with separate accumulators) was used for each circle radius (quarters vs. penny).

Hough transform: pros and cons

Pros

- All points are processed independently, so can cope with occlusion, gaps
- Some robustness to noise: noise points unlikely to contribute *consistently* to any single bin
- Can detect multiple instances of a model in a single pass

Cons

- Complexity of search time increases exponentially with the number of model parameters
- Non-target shapes can produce spurious peaks in parameter space
- Quantization: can be tricky to pick a good grid size